

Simpson's Paradox — Corpus Summary

A digital mirror of Simpson's Paradox — the paradox falls out of the DAG as a derived fact, never modeled directly.

Generated 2026-06-29 from live vw_* views — same data as /conclusions. SSoT is the rulebook JSON.

Scope boundary (conc-09)

The instrument classifies allocation geometry only. Causal interpretation is external, gated by AdjustmentAppropriate. Berkeley: highest AllocationDistortion (0.193) but CausalRole=contested — geometry and confounding are not the same thing.

Scope & limits

Theorem conclusions follow from definitions (regression checks for the transpiler). Domain conclusions and corpus discovery findings are statistics on a curated convenience sample — directional, not inferential. See conc-14 before treating cross-domain patterns as laws.

Allocation sweep: 10 steps per study, $f \in [0.05, 0.95]$ (SweepStudyConfig). Latent flip = Type D at observed allocation + PooledGapCrossesZero somewhere in that range — not necessarily at observed allocation.

Epistemic summary (tiered)

Proved by construction (theorem): 1

Established in instrument (instrument · taxonomy · methodology · scope): 13

Observed in this corpus (domain — provisional): 3

Total conclusions: 17

Corpus hypotheses (loop-61): 4 / 4 confirmed

Consistency checks (definition-linked): 1 / 1 pass

Latent Type-D (under sweep contract above): 45 / 64 (70.3%)

Studies in corpus: 96

Distortion taxonomy: A 7 · B 11 · C+ 8 · C" 6 · D 64

Pre-registered discovery (loop-61)

Consistency checks — confirm implementation matches stated algebra (same tier as conc-12 / inv-signal-purity-sign-flip). Not independent discoveries about nature.

H-purity — PASS (consistency check)

Every sign-flip study has SignalPurity strictly below 0.5.

Expected: SignFlipSignalPurityMax < 0.5

Observed: maxPurity=0.483

Corpus hypotheses — contingent patterns on this convenience sample. Directional inequalities only; not inferential. See conc-14.

H-domain-dist — PASS (corpus hypothesis)

Epidemiology studies carry higher mean AllocationDistortion than education studies.

Expected: EpidemiologyAvgDistortion > EducationAvgDistortion

Observed: epi=0.049 edu=0.004

H-econ-zero — PASS (corpus hypothesis)

Economics-domain studies produce zero sign-flips in 2×K encoding.

Expected: EconomicsSignFlipCount = 0

Observed: flips=0

H-latent-d — PASS (corpus hypothesis)

More than half of Type-D (SAFE) studies would flip their pooled winner under counterfactual allocation reweighting.

Expected: LatentTypeDFraction > 0.5

Observed: fraction=0.703

H-small-effect — PASS (corpus hypothesis)

Allocation-stable Type-D studies carry larger observed pooled gaps than latent Type-D studies.

Expected: AvgPooledGapStableD > AvgPooledGapLatentD

Observed: stable=0.102 latent=0.022

Proved (by construction) (1)

conc-12-signal-purity-theorem · Proved (by construction)

SignalPurity < 0.5 is a necessary condition for allocation-driven reversal

Witnessed across all reversal studies (type A + B) in the 90+ study corpus: every study with AllocationDirection='reversal' has SignalPurity < 0.5. The algebraic proof: reversal requires SignedPooledGap and CorrectedGap to have opposite signs, which requires AllocationDistortion > |CorrectedGap|, which is exactly SignalPurity < 0.5. inv-signal-purity-sign-flip enforces this as a DAG invariant. AvgSignalPurityReversal "H 0.35 vs AvgSignalPurityNonReversal "H 0.65 — gap of ~0.30.

Loop replay: [loop-44 \(5a14dcb. Sun Jun 28 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Established in instrument (8)

conc-03-binary-excludes-partial · Established in instrument

The unanimity criterion (IsReversal) incorrectly excludes Berkeley 1973 despite its being a Simpson-type paradox

Berkeley: WSGSUM=" 0.0642, SignedPooledGap=+0.1292, IsSignFlip=TRUE, AllocationDistortion=0.1934 — greater distortion than kidney-1986. But IsReversal=FALSE because dept-C marginally favors male. The unanimity criterion is a strict subtype, not the primary definition.

Loop replay: [loop-17 \(ca31095. Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Protecting invariants: 4

conc-04-type-c-exists · Established in instrument

Allocation can compress the pooled signal without flipping it (type-C): a class invisible to IsReversal and IsSignFlip but visible to AllocationDistortion

compressed-synthetic: A wins every stratum and wins pooled, but the pooled margin is 6pp rather than the equal-weight 10pp. AllocationDistortion=0.0400, IsSignFlip=FALSE. The pooled winner is correct; the effect size is wrong by 4pp.

Loop replay: [loop-19 \(ca31095. Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Protecting invariants: 4

conc-05-real-data-generalizes · Established in instrument

Two real published studies (reintjes-2000, radelet-1981) hydrate as type-A full reversals without schema changes

Both studies entered via 05b-customize-data.sql. Witnessed: reintjes AllocationDistortion=0.0516, DistortionType=A; radelet AllocationDistortion=0.0684, DistortionType=A. No DAG or schema change was needed — the hydration contract generalizes.

Loop replay: [loop-20b \(ca31095. Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

conc-06-policy-derivable · Established in instrument

PolicyImplication is domain-neutral: the same formula produces researcher actions across medicine, law, and education

All five current studies (kidney-1986, berkeley-1973, compressed-synthetic, kidney-balanced, balanced-synthetic) plus two real studies produce unambiguous PolicyImplication values from DistortionType alone. No domain-specific knowledge is encoded in the formula.

Loop replay: [loop-21 \(ca31095, Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

conc-07-model-self-portrait · Established in instrument

ModelSummary witnesses the instrument's own epistemic coverage as a first-class derived fact

ModelSummary aggregates ReversalCount, ExplainedCount, TypeA/B/C/D counts, and AvgAllocationDistortion across all studies. The model describes its own completeness in a single row.

Loop replay: [loop-10 \(b96e0fe, Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

conc-11-reversal-recovery · Established in instrument

The model can produce the allocation-corrected winner directly from existing DAG fields — no new data required

CorrectedGap = WeightedStratumGapSum (already derived in loop-12/17). CorrectedWinner is a sign-read on CorrectedGap.

CorrectedVsPooledAgreement is TRUE for Type-D studies and FALSE for Type-A/B studies. For kidney-1986:

CorrectedGap=+0.0537, CorrectedWinner=A, CorrectedVsPooledAgreement=FALSE (pooled said B; corrected says A). The reversal recovery is free — it costs zero new data.

Loop replay: [loop-27 \(ca31095, Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Protecting invariants: 3

conc-15-latent-flip-potential · Established in instrument

Type-D at observed allocation does not imply allocation-stable — 70% of SAFE studies flip under counterfactual reweighting

Sensitivity over sweep range $f \in [0.05, 0.95]$; not a claim about realistic reallocation alone.

LatentTypeDCount=45, TypeDCount=64, LatentTypeDFraction=0.703. LatentFlipPotential=TRUE when DistortionType=D AND PooledGapCrossesZero=TRUE. Examples: uc-systemwide-admissions (Type D, sweep range 0.56), recsys-offline-eval (Type D, sweep range 0.56). The ScreeningTier SAFE label describes observed allocation only; sweep reveals latent flip potential.

Loop replay: [loop-61 \(7b3938e, Mon Jun 29 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Protecting invariants: 2

conc-16-small-effect-fragile · Established in instrument

In this corpus, allocation-stable Type-D rows carry larger observed pooled gaps than latent Type-D rows

AvgPooledGapStableD=0.102 across 19 stable Type-D studies; AvgPooledGapLatentD=0.022 across 45 latent Type-D studies. Stable studies also have lower sweep range (avg 0.072 vs 0.211). Corpus-descriptive pattern on loop-61 convenience sample — not a deductive law. Small observed effects tend to be allocation-fragile even when currently classified SAFE.

Loop replay: [loop-61 \(7b3938e, Mon Jun 29 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Taxonomy (1)

conc-13-cplus-cminus-opposite-errors · Taxonomy

C+ (amplification) and C- (compression) are epistemically opposite errors with opposite clinical implications

titanic-1912 (C+, DistortionRatio=1.359): allocation inflates first-class survival advantage from 27.5pp corrected to 37.4pp pooled — overconfidence. melanoma-altman-1991 (C+, DistortionRatio=2.087): allocation doubles apparent female survival advantage — severe overconfidence. hannan-1994 (C-, DistortionRatio=0.314): allocation compresses CABG advantage from 8.9pp corrected to 2.8pp pooled — underconfidence, risk of premature treatment abandonment. The prior type-C category hid this distinction. Now separated b...

Loop replay: [loop-43 \(5a14dcb, Sun Jun 28 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Methodology (2)

conc-10-instrument-specifiable · Methodology

The full instrument can be specified precisely enough for any researcher to apply it to a new study

Witnessed in loop-26: InstrumentSpec table enumerates 5 input fields, 10 derived coordinates, 2 classification rules, and 3 adapter-requirement rows. Machine-verifiable against the full 90+ study corpus.

Loop replay: [loop-26 \(ca31095, Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

conc-14-scale-is-discovery-path · Methodology

Cross-domain SignalPurity correlations require a pre-registered unbiased corpus of 50+ studies to distinguish pattern from selection artifact

Partial domain patterns witnessed across 90+ studies (loop-61): epidemiology avg distortion 0.050 vs education 0.004; legal/sports/epidemiology flip rates ~27% vs economics 0%. Full causal domain! DistortionType mapping still requires pre-registered expansion beyond convenience sample, but the instrument now supports DiscoveryHypotheses/Findings for ongoing corpus-scale tests.

Loop replay: [loop-61 \(7b3938e, Mon Jun 29 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Scope / open question (2)

conc-08-type-c-real · Scope / open question

Does the published literature contain real type-C studies (compression without sign flip)?

Witnessed in loop-24: hannan-1994 (CABG surgery), titanic-1912, melanoma-altman-1991, and coffee-tverdal-2020 all hydrate as DistortionType=C. AllocationDistortion>0.01, IsSignFlip=FALSE — the correct pooled direction but compressed magnitude. The type-C cell is not a synthetic pathology.

Loop replay: [loop-24 \(ca31095, Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

conc-09-causal-vs-geometric · Scope / open question

AllocationDistortion is a geometric measure; causal interpretation requires external domain knowledge, gated by AdjustmentAppropriate

Berkeley has the highest AllocationDistortion (0.1934) but CausalRole=contested (mediator vs confounder), proving that geometric severity and causal confounding are not the same thing. The instrument is deliberately geometric: it classifies distortion type and flags the corrected winner from allocation arithmetic alone. Whether it is safe to act on that correction as a causal claim is answered by AdjustmentAppropriate — which gates on ConditioningRisk and CausalClaimStatus. The two layers are complementary, not ...

Loop replay: [loop-30 \(ca31095, Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Protecting invariants: 1

Observed in this corpus (3)

conc-01-paradox-is-derived · Observed in this corpus

Simpson's Paradox is a derived fact from the DAG, not a modeled entity

IsReversal (and later IsSignFlip) falls out as a calculated field on TreatmentRankings, derived from aggregated CaseCell counts. No 'ReversalDetection' entity was ever needed.

Loop replay: [loop-04 \(b96e0fe, Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Protecting invariants: 2

conc-02-reversal-from-allocation · Observed in this corpus

The reversal in kidney-1986 is entirely an allocation artifact — same rates with balanced allocation produce no reversal

kidney-balanced uses identical stratum success rates to kidney-1986 but with equal allocation (175 cases per cell). IsReversal=FALSE, AllocationDistortion=0. The paradox is algebraically isolated as a property of allocation, not of treatment efficacy.

Loop replay: [loop-14 \(49f01e7, Sat Jun 27 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Protecting invariants: 3

conc-17-economics-flip-free · Observed in this corpus

Economics-domain 2×K encodings in this corpus produce zero sign-flips — composition studies read as Type D or C"

This corpus only (n=8 economics studies); see conc-14 for inferential limits.

EconomicsSignFlipCount=0 across 8 economics studies (6 Type D, 2 Type C"). Contrast epidemiology (6 sign-flips / 22 studies, 27.3%). Composition and wage-decomposition paradoxes compress or stabilize rather than flip in this instrument's 2×K encoding.

Loop replay: [loop-61 \(7b3938e, Mon Jun 29 2026 00:00:00 GMT-0500 \(Central Daylight Time\)\)](#)

Invariant health

Critical invariant failures: [0](#)

inv-pooled-gap-wanders: Pooled gap should move under allocation reweighting for reversal studies. Warning-only: small-magnitude reversals may have narrow sweep range. (fail=1)

Most misleading studies (by paradox strength)

COVID-19 CFR: Italy vs China by Age (von Kügelgen 2021) — type C+, strength 0.2743, distortion 0.529, purity 3.6% (CAUTION)

Triathlon On-Time Finish: Sex by Age Cohort — type B, strength 0.1455, distortion 0.245, purity 28.9% (DANGER)

Smoking and Mortality by Age (Cochran / Surgeon General 1964) — type B, strength 0.1091, distortion 0.167, purity 25.7% (DANGER)

UC Berkeley Graduate Admissions (Bickel et al. 1975) — contested confounder — type B, strength 0.0969, distortion 0.193, purity 24.9% (DANGER)

UC Berkeley Admissions: Six Largest Departments (Bickel et al. 1975) — type B, strength 0.0944, distortion 0.184, purity 18.8% (DANGER)

Exercise and Cholesterol Control by Age Group (Framingham-style) — type D, strength 0.0800, distortion 0.008, purity 95.2% (SAFE)

Baseball Batting Averages: Lefebvre vs Fairly World Series (Day 1994 / baseball paradox canon) — type B, strength 0.0606, distortion 0.082, purity 20.9% (DANGER)

Whickham Smoking Cohort — 20-year Mortality by Smoking Status (Appleton et al. 1996) — type B, strength 0.0538, distortion 0.119, purity 26.9% (DANGER)

Rulebook & repo (full derivation)

Formulas, Loops build history, InvariantChecks, and Conclusions rows live in the rulebook JSON. README.md carries the full framing.

Project folder: <https://github.com/EffortlessAPI/effortless-rulebooks/tree/main/rulebook-examples/simpsons-paradox>

README (full repo guide): [README.md](#)

Rulebook SSoT (formulas, Loops, InvariantChecks, Conclusions): [effortless-rulebook/simpsons-paradox-rulebook.json](#)

Rulespeak (plain-English business rules): [rulespeak/rulespeak.md](#)

Postgres substrate (views, init-db.sh): [effortless-postgres/README.md](#)

Standalone HTML explorer: [simpsons-paradox-explorer.html](#)

Live API (with ./start.sh): GET localhost:3001/api/invariant-checks · /api/treatment-rankings · /api/model-summary

Loop landing tags (replay checkpoints):

sp-loop-milestone-01-loops-01-10: [conc-01-paradox-is-derived](#)
sp-loop-milestone-03-loops-11-15: [conc-02-reversal-from-allocation](#)
sp-loop-milestone-04-loops-16-31: [conc-03-binary-excludes-partial](#)
sp-loop-milestone-04-loops-16-31: [conc-04-type-c-exists](#)
sp-loop-milestone-04-loops-16-31: [conc-05-real-data-generalizes](#)
sp-loop-milestone-04-loops-16-31: [conc-06-policy-derivable](#)
sp-loop-milestone-01-loops-01-10: [conc-07-model-self-portrait](#)
sp-loop-milestone-04-loops-16-31: [conc-08-type-c-real](#)
sp-loop-milestone-04-loops-16-31: [conc-09-causal-vs-geometric](#)
sp-loop-milestone-04-loops-16-31: [conc-10-instrument-specifiable](#)